



Summer Project

Breaking Ice Into Business: A Traction to Subscription Model

Endterm Report

Abstract

This documentation of the end-term-evaluation provides the details of concepts taught in this project.

Contents

1	Introduction to consulting	4
1.1	What is Consulting	4
1.2	Types of Consulting Services	4
1.3	What is the work of a consultant?	4
1.4	Common Misconceptions	5
2	What are Consulting Frameworks?	5
2.1	SWOT Analysis	6
2.2	PESTEL Analysis	6
2.3	Porter's Five Forces	7
2.4	Ansoff Growth Matrix	8
3	Market Sizing Analysis	9
3.1	Why are TAM, SAM and SOM important ?	9
3.2	Calculation Methods	9
3.2.1	Top-Down Approach	10
3.2.2	Bottom-up Approach	10
3.3	TAM (Total Addressable Market)	10
3.4	SAM (Serviceable Addressable Market)	10
3.5	SOM (Serviceable Obtainable Market)	11
4	Subscription Pricing Models	11
4.1	Flat-rate	11
4.2	Tiered	11
4.3	Usage-Based	12
4.4	Per-Added-Module	12
4.5	Per-User	12
5	Subscription Business Model	12
5.1	What is a Subscription Business Model?	12
5.2	Advantages of the Subscription Business Model	12
5.3	Popular Subscription Business Models	13
5.3.1	SaaS (Software-as-a-Service)	13
5.3.2	Subscription Boxes	13
5.3.3	Publications & Newsletters	13
5.3.4	Media Subscriptions	13
5.4	Metrics all Subscription Models track	13
5.4.1	Monthly Recurring Revenue (MRR)	13
5.4.2	Annual Recurring Revenue (ARR)	13
5.4.3	Average Revenue Per User (ARPU)	13
5.4.4	Customer Lifetime Value (CLTV)	14
5.4.5	Customer Acquisition Cost (CAC)	14
5.4.6	Churn Rate	14
5.5	Subscription Box Financial Model	14
5.6	Excel-Based Subscription Financial Model	14
6	Cost-Benefit Analysis	15
6.1	What is Cost-Benefit Analysis ?	15
6.2	Time Factor in Cost-Benefit Analysis	16
6.3	Best Practices in Cost-Benefit Analysis	16
7	How to Calculate the Break-Even Point	17
7.1	What is the Break-Even Point ?	17
7.2	How to Calculate the Break-Even Point ?	17
7.2.1	In Units	17
7.2.2	In Revenue	18

8	Metrics and KPI	18
8.1	Customer Retention Rate (CRR)	18
8.2	Net Promoter Score (NPS)	18
8.3	Churn Rate	18
8.4	Monthly Active Users (MAU)	18
8.5	Daily Active Users (DAU)	18
8.6	Activation Rate	19
8.7	Average Session Length	19
8.8	Customer Acquisition Cost (CAC)	19
8.9	Lifetime Value (LTV)	19
8.10	LTV : CAC Ratio	19
8.11	Story Completion Rate	19
8.12	Engagement Rate	19
8.13	Virality Coefficient	19
8.14	Revenue Metrics	19
8.15	Funnel Conversion Rates	19
8.16	Bounce Rate	20
8.17	DAU/MAU Ratio	20
9	Introduction to Python Libraries	20
9.1	NumPy: Numerical Python	20
9.1.1	Overview	20
9.1.2	Key Features and Functionalities	20
9.1.3	Example Usage	20
9.2	Pandas: Data Analysis Library	21
9.2.1	Overview	21
9.2.2	Key Data Structures	21
9.2.3	Key Features and Functionalities	21
9.2.4	Example Usage	21
9.3	Matplotlib: Plotting Library	21
9.3.1	Overview	21
9.3.2	Example Usage	21
10	Machine Learning	22
10.1	Introduction to ML	22
10.2	Mean based model	22
11	Linear Regression	22
11.1	Introduction	22
11.2	Simple Linear Regression	23
11.3	Multiple Linear Regression	23
11.4	Assumptions	23
11.5	Mathematical Formulation	23
11.6	Example and Interpretation	23
11.7	Applications	24
11.8	Linear regression and gradient descent	24
11.9	Non iterative approach	25
12	Unsupervised Learning	25
12.1	How Unsupervised Learning Works?	25
12.2	Clustering in Unsupervised Learning	26
12.3	How Clustering Reveals Structure	26
12.4	K-Means Clustering	26
12.4.1	How K-Means Clustering Works	26
12.4.2	Intuition Behind K-Means Clustering	27
12.4.3	WCSS and the Elbow Method	27
12.4.4	Random Initialization Trap in K-Means	28
12.5	Hierarchical Clustering	28
12.5.1	What is Hierarchical Clustering?	28

12.5.2	Dendrogram: A Tree for Clusters	28
12.5.3	Types of Hierarchical Clustering	29
12.5.4	Inter-Cluster Distance (Linkage Methods	29
12.5.5	How to Select the Number of Clusters?	29
13	Principal Component Analysis (PCA)	30
13.1	What is PCA?	30
13.2	Benefits of PCA	30
13.3	Why Variance is Important?	30
13.4	Covariance	30
13.5	Geometric Interpretation of PCA	31
14	Decision Tree	31
14.1	What is a Decision Tree?	31
14.2	How Does a Decision Tree Work?	32
14.3	Decision Tree Structure Overview	32
14.4	Splitting Criteria in Classification Trees	32
14.5	Information Gain	32
14.5.1	Formula:	32
14.6	Decision Trees for Regression	32

CONSULT

1 Introduction to consulting

1.1 What is Consulting

Consulting is a professional service where experts (consultants) provide **objective advice, insights and strategic guidance** to businesses or individuals facing complex problems or needing support in achieving goals. It's not just advice, it often involves analysis, solution design, and sometimes execution support.

The purpose of consulting is to offer objective insights, strategic recommendations, and practical assistance to help clients overcome obstacles, capitalize on opportunities, and enhance their overall performance. Consultants bring a fresh perspective and often employ a systematic and analytical approach to assess the client's situation, identify areas for improvement, and develop tailored strategies and action plans.

Consulting services generally include:

- Marketing and sales
- Financial advisory services
- IT and technology consulting
- Business strategy and management
- Operations and supply chain management

The Growth of the Consulting Industry

- In the U.S. alone, the consulting industry is valued at around **\$55 billion**.
- The **overall industry** grew by **7.7%**, while **management consulting** grew by about **2%**.
- There are **over 900,000 consulting firms** in the U.S.—ranging from large players like McKinsey or Deloitte to solo consultants.

This growth is driven by:

- Increasing complexity in business and technology
- The need for specialized skills and outside perspectives
- Cost-efficiency for companies to hire short-term experts

Their role is both **strategic (planning)** and **tactical (execution)** depending on the project.

1.2 Types of Consulting Services

Consulting spans many specialties. Common ones include:

- **Strategy Consulting:** Long-term planning, market entry, competitive analysis.
- **Operations Consulting:** Streamline processes, increase efficiency, reduce costs.
- **Financial Consulting:** Budgeting, forecasting, investments, risk management.
- **Marketing Consulting:** Branding, customer acquisition, digital strategy.
- **IT Consulting:** Technology selection, system implementation, cybersecurity.
- **HR/People Consulting:** Talent acquisition, organizational design, DEI practices.

1.3 What is the work of a consultant?

The role of a consultant can be complex and varied depending on their area of expertise. However, in general, they work closely with company stakeholders, conduct research and analysis, provide recommendations, and assist with the implementation of solutions. They do this by bringing an external perspective, industry knowledge, and best practices to help organizations achieve their goals and improve overall performance.

The key responsibilities of a consultant typically include:

-
- **Analysis and assessment:** Consultants thoroughly analyze the client's business operations, processes, strategies, and challenges. They gather data, conduct research, and perform diagnostic assessments to gain a deep understanding of the client's current situation.
 - **Problem-solving and strategy development:** Based on their analysis, consultants identify problems, bottlenecks, and areas for improvement within the client's organization.
 - **Process improvement and optimization:** Consultants work with the client's team to optimize business processes, streamline operations, and enhance efficiency.
 - **Change management:** Consultants help clients navigate transitions, restructures, mergers, or acquisitions. They do this by guiding communication, stakeholder engagement, and managing resistance to change.
 - **Organizational design and restructuring:** Consultants assist in designing or restructuring organizational structures, roles, and responsibilities. They evaluate the client's current organizational design, propose changes, and provide guidance on implementing new structures.
 - **Strategic planning and execution:** Consultants help define strategic priorities, identify growth opportunities, and conduct market analysis. They also create implementation roadmaps to ensure successful execution.
 - **Financial analysis and planning:** Consultants assess the client's financial health and identify areas of improvement. Then, they provide recommendations for maximizing profitability and managing financial risks.
 - **Training and skill development:** Consultants may deliver training programs in areas such as leadership development, project management, sales, or process improvement.
 - **Performance measurement and evaluation:** Consultants establish performance metrics and evaluate outcomes and identify any gaps. They also provide feedback to help the client continuously improve and achieve sustainable results.
 - **Knowledge transfer and client empowerment:** A fundamental responsibility of consultants is to share expertise, best practices, and insights gained. This is so the client can sustain improvements independently and build internal capabilities.

1.4 Common Misconceptions

Misconception 1: "Consultants are only needed when things are broken."

Truth: They are often hired to **optimize**, **scale**, or **innovate**, not just fix problems.

Misconception 2: "Consultants are generalists."

Truth: Most are **highly specialized** in a field (e.g., data analytics, healthcare strategy, SaaS pricing).

Misconception 3: "Hiring a consultant is too expensive"

Truth: The ROI can be significant if the right consultant is hired for the right problem.

2 What are Consulting Frameworks?

Consulting frameworks are **structured thinking tools** that consultants use to **break down complex business problems** into manageable parts. These frameworks offer a **systematic approach** to analyze issues, develop strategies, and make data-driven decisions. They are widely used in **management consulting**, **case study interviews**, and **strategic planning** for businesses in all industries.

There are multiple consulting frameworks but we will see a few of them:

1. SWOT Matrix
2. Pestel analysis
3. Porter's Five Forces
4. Ansoff's Growth Matrix

2.1 SWOT Analysis

SWOT Analysis is more of a mini-framework, specifically for quickly evaluating a single company in an industry. It evaluates both intrinsic and extrinsic factors based on facts and data to create realistic predictions.

SWOT is effectively a quick, high-level market landscape/competitive dynamics analysis arranged using the following terminology:

- **Strengths:** Company strengths within an industry.
- **Weaknesses:** Company weaknesses within an industry
- **Opportunities:** Company opportunities available within the industry (or potentially by branching into a new industry).
- **Threats:** Company threats within the industry (or potentially from companies whose primary business is in another, related industry, or from disruptive technologies that potentially threaten all companies in an industry)

For Example, let us take the case study of Zomato:

- **Strengths**
 1. Effective Marketing Strategy
 2. Strong brand recognition
- **Weaknesses**
 1. Low customer retention
 2. High platform fee
- **Opportunities**
 1. Tapping into remote locations
 2. Market growth(by CAGR)
- **Threats**
 1. Data security issues
 2. Legal threats

2.2 PESTEL Analysis

PESTEL analysis is a method that managers can use to evaluate any major external factors that may influence their operations to help them stay competitive in the market. Those using this framework can maximize their current conditions and prepare for future changes that may bring opportunities or challenges.

- **P – Political:**The political analysis examines how government regulations and other factors, such as tax guidelines, could affect trading and business.
- **E – Economic:**The economic analysis reviews financial issues, such as inflation and interest rates, that may affect an organization's success.
- **S – Social:**The social analysis evaluates characteristics, such as lifestyle and education, of their customers to understand their purchase needs better.
- **T – Technological:**The technological analysis determines how technology factors, such as technological advances, can positively or negatively influence how customers respond to a new product or service.
- **E – Environmental:**This covers ecological and environmental aspects such as climate change, environmental regulations, and sustainability efforts.
- **L – Legal:**It includes laws regarding consumer protection, antitrust laws, employment laws, and health and safety regulations

Let's do a PESTEL analysis on whether Tesla should enter the Indian market or not:

1. **Political**

-
- (a) The government is supporting local players which is a setback for Tesla.
 - (b) Increased custom duties(200%) in India cause resistance for international players to enter the Indian market.
2. **Economical**
- (a) In India targeting to launch a car under 20 lakhs will be a good choice.
 - (b) But, Tesla has very high R&D costs per car(approx 3000\$ /car)
3. **Social**
- (a) A very high proportion of the market acquired by TATA EVs
 - (b) People usually have resistance to switching their vehicle preferences from petrol to EVs.
4. **Technological**
- (a) Ground clearance issues are prominent in India.
 - (b) Safety concerns about self-driving features can pose a threat to Tesla.
 - (c) Tesla has to get better energy density and faster charging time in their batteries to lure Indian customers.
5. **Environmental**
- (a) The adverse effects of high temperatures on batteries can potentially threaten the life of Tesla cars in India.
 - (b) High water consumption in manufacturing is also a big factor to look at.
6. **Legal**
- (a) Navigating the GST framework and ensuring compliance with tax regulations could be challenging.
 - (b) Acquiring land for manufacturing plants can be a complex process due to India's land use regulations and bureaucratic procedures

2.3 Porter's Five Forces

Porter's five forces are used to identify and analyze an industry's competitive forces. The model guides businesses in determining the intensity of competition and potential profitability within their market, helping them better understand where power lies in their sector.

- **Competitive Rivalry:** This force examines the degree of competition between existing firms in the industry. High rivalry limits profitability as competitors engage in price wars, advertising battles, and new product launches
- **Factors to Consider:** Number of competitors, rate of industry growth, product or service differentiation, brand loyalty, barriers to exit and switching costs..
- **Threat of New Entrants:** This force assesses how easily new competitors can enter the market. High threat of new entrants can decrease market share and profitability for existing companies
- **Factors to Consider:** Barriers to entry (such as economies of scale, capital requirements, access to distribution channels, and government regulations), brand recognition, and customer loyalty, capital requirements, government policies, access to distribution channels, switching costs.
- **Bargaining Power of Suppliers:** This force analyzes the power that suppliers have over the prices and quality of materials or services they provide. High bargaining power of suppliers can increase costs and reduce profitability for firms in the industry.
- **Factors to Consider:** Number and size of suppliers, uniqueness of supplier products or services, switching costs to change suppliers, focal company's ability to substitute.
- **Bargaining Power of Buyers:** This force examines the influence that customers have on prices and quality. High bargaining power of buyers can drive prices down and demand higher quality or more services, thus reducing profitability.
- **Factors to Consider:** Number of buyers, volume of purchases, differentiation of products or services, price sensitivity, buyer's information and switching costs to change suppliers.

- **Threat of Substitutes:** This force evaluates the likelihood that customers will switch to different products or services that perform the same function. High threat of substitutes can limit the potential for price increases and reduce industry profitability.
- **Factors to Consider:** Availability of substitute products or services, performance and cost of substitutes, switching costs and buyer willingness to switch to substitutes.

2.4 Ansoff Growth Matrix

The Ansoff Growth Matrix is a strategic planning tool designed to help businesses identify and plan their growth strategies. It focuses on two dimensions: products and markets. The matrix is divided into four quadrants, each representing a different growth strategy:

- **Market Penetration:** Selling more existing products to existing markets.
- **Market Development:** Entering new markets with existing products.
- **Product Development:** Creating new products for existing markets.
- **Diversification:** New products in new markets (most risky).



1. Market Penetration: Increasing market share for existing products in existing markets its objective is to grow market share within the current market by increasing the usage of current products among existing customers and attracting competitors' customers.

Strategies:

- **Increase Usage:** Encourage existing customers to use the product more frequently or in larger quantities.
- **Attract Competitors' Customers:** Offer better services, lower prices, or improved product features to lure customers away from competitors.
- **Enhance Distribution:** Expand the availability of the product by improving distribution channels.
- **Promotions and Marketing:** Use aggressive advertising, promotions, and sales tactics to boost market presence.

2. Market Development: Entering new markets with existing product, its objective is to grow by tapping into new geographic regions, new customer segments, or new uses for existing products.

Strategies:

- **Geographic Expansion:** Entering new regions or countries where the product is not currently sold.
- **Target New Segments:** Identifying and targeting new customer demographics that were previously untapped.
- **New Distribution Channels:** Using different distribution methods to reach new customers (e.g., online sales, direct sales).

3. Product Development: Developing new products for existing markets. Its objective is to grow by innovating and offering new products or improving existing products to the current customer base.

Strategies:

- **Product Innovation:** Creating new features or entirely new products that appeal to the existing market.
- **Product Improvement:** Enhancing the features, quality, or performance of existing products.
- **Extending Product Lines:** Adding new products that complement or are related to the existing product line.

4. Diversification: Entering new markets with new products its objective is to grow by developing new products for new markets, which can be related or unrelated to the current business.

Strategies:

- **Acquisitions and Mergers:** Acquiring or merging with companies in new industries or markets.
- **New Ventures:** Creating new businesses or subsidiaries to enter new markets.
- **Strategic Alliances:** Forming partnerships or alliances with companies in different industries.


3 Market Sizing Analysis

Market sizing is a crucial step in understanding the potential of a business or product within a given market. It involves estimating how large the opportunity is and how much of it can realistically be captured. The three key components used in market sizing are TAM, SAM, and SOM.

3.1 Why are TAM, SAM and SOM important ?

Several reason why calculating TAM, SAM and SOM are beneficial to various businesses are -

1. **Market potential assessment:** These metrics help businesses estimate the potential market size for their products or services.
2. **Investment decisions:** Investors and venture capitalists often use TAM, SAM, and SOM to evaluate the potential of a startup or a new project before making investment decisions.
3. **Product development and market entry strategies:** Knowing the market size can help businesses decide whether to develop new products or enter new markets.
4. **Competitive analysis:** These metrics allow businesses understand their competitive position in the market and identify areas where they can expand.



3.2 Calculation Methods

There are two primary approaches to calculating TAM, SAM, and SOM which are top-down and bottom-up.

3.2.1 Top-Down Approach

The top-down begins with a broad market overview and narrows to the target market segments. In this method, we typically use market research data to estimate the total market size and then apply various assumptions and filters to determine the SAM and SOM.

Advantages :

1. Provides a broad market perspective.
2. Is relatively quick and easy to execute.

Limitations :

1. May rely heavily on secondary data, which can be outdated or inaccurate.
2. Can be overly optimistic and can lead to an overestimation of the potential market size.

3.2.2 Bottom-up Approach

The bottom-up approach starts with the small components and builds up to the whole system. It starts with a detailed analysis of the company's target market segments and builds up to estimate the TAM, SAM, and SOM.

Advantages :

1. Provides a more accurate and detailed market size estimation.
2. Is more reliable for setting realistic revenue targets and sales forecasts.

Limitations :

1. Can be time-consuming and resource-intensive.
2. Harder to get investor attention without showing larger potential.

3.3 TAM (Total Addressable Market)

TAM represents the total revenue opportunity available for a product or service in a specific market or industry. It is the maximum revenue a company could generate if it held 100% market share.

How to calculate TAM ?

Top-down approach :

We start with using broader market data to estimate the market size and then refine the TAM estimate based on the company's target market segments. While this approach requires access to accurate, market-level data, the data collection is the hard part.

Bottom-up approach :

We start with doing the data collection to estimate the potential market size for the company's target market segments. We can aggregate the estimates for each segment to determine the overall TAM. This method is much more labour intensive, but in the absence of accurate market data, it might be our only good option.

3.4 SAM (Serviceable Addressable Market)

SAM is the portion of TAM that a company can realistically serve with its current products and distribution channels. Considering the company's limitations and competition, it is a more focused estimate of the potential market size.

How to calculate SAM ?

Top-down approach:

We start with the TAM estimate and apply various filters and assumptions to refine the SAM estimate based on the company's specific target market segments.

Bottom-up approach:

Similar to the TAM approach, we collect primary data, like sales data, and use this information to estimate the potential market size for each target market segment. Then, we could add the estimates for each segment together to determine the overall SAM.

3.5 SOM (Serviceable Obtainable Market)

SOM is the market share a company can realistically achieve within its SAM in a specific timeframe. It is the most conservative estimate of the potential market size and helps businesses set realistic revenue targets.

How to calculate SOM?

Calculating SOM involves generating a reasonable estimate of the portion of your SAM that your company can achieve in a certain period. It is a bit trickier to calculate SOM than the other two, and a simple bottom-up or top-down approach won't necessarily serve us well.

However, we use these tips to estimate SOM very well -

1. Understanding Competitive Edge: We start by analysing what makes our product or service stand out.
2. Market Share Analysis: Looking at our current market share in the SAM and assessing how much we can grow based on factors like marketing efforts, product development, and sales strategies.
3. Assessing Limitations and Barriers: Identifying any barriers to growth, such as limited production capacity, budget constraints, or strong competition.
4. Setting Realistic Goals: Based on these analyses, we ultimately set a realistic percentage of the SAM we believe our company can obtain which should be a specific, time-bound target.

4 Subscription Pricing Models

Subscription pricing is a revenue model where customers pay a recurring fee, typically monthly or annually, for access to a product or service, such as cloud-based storage and multimedia content.

4.1 Flat-rate

Flat-rate, also known as fixed pricing, offers users a single price for all features of the offering. This model is best for companies with a product that has limited features and a single buyer persona. Flat-rate does not work well for companies in which resource costs might vary significantly from user to user.

Advantages :

1. Simple to understand, communicate and sell
2. Due to its one-size-fits-all approach, it frees up time for companies to focus on monetization, acquisition and retention instead of tailoring pricing strategy

4.2 Tiered

This model offers packages with various features and product combinations at various price points. This allows sellers to segment the prices of their products and services based on specified target markets. Tiers are generally designated as basic, standard and premium.

This model is best for companies that may have many product features and a diverse customer base with varying needs and budgets.

Advantages :

1. Caters to multiple buyer personas by offering multiple price points.
2. Maximizes the lifetime value of the customer, as they have the option to upgrade/downgrade as needs change.

4.3 Usage-Based

The usage-based model is also referred to as a consumption model and pay-as-you-go. It directly relates the cost of a product to its level of consumption, typically involving a base rate with an additional usage rate.

This model is best for products or services of which customers' usage is likely to vary widely.

Advantages :

1. Offers customers the most flexibility and attracts new customers with the low initial costs because of low usage.
2. Heavy users that may require more resources are charged more than less-frequent users.

4.4 Per-Added-Module

This model, product is priced on the basis of the functionality offered to the customers. There is a "base product" and the option to add modules for more functionality — at a higher cost.

This model is best for companies which provide modular functionality that can easily be added to the core product.

Advantages :

1. Lower Entry Barrier
2. Scalable Revenue for the Provider

4.5 Per-User

The per user pricing model is a subscription-based pricing model where the cost of a product or service is determined by the number of users who can access it.

This model is best for products that involves multiple users, particularly those that demand teamwork or collaboration.

Advantages :

1. Pricing scales linearly with the number of users, making it easy to estimate costs.
2. Each user gets their own login and usage can be tracked individually.

5 Subscription Business Model

5.1 What is a Subscription Business Model?

When a business charges a recurring fee for its services and products, it follows a subscription model. The recurring charges can be on an annual or monthly basis.

What makes the subscription business model so popular for customers is the convenience and personalization that it offers, and for companies, it is the predictability and sustainability.

5.2 Advantages of the Subscription Business Model

1. **Predictable Revenue** : The subscription model creates a predictable future revenue stream. It also helps in keeping the business operations up and running in times of turbulence and financial fluctuations.
2. **Better Customer Relationships** : Customers often opt for subscriptions because of the convenience factor that comes with options to personalize the product. Customized subscriptions help develop customer loyalty.
3. **Reduction of Retention costs** : Since subscribers automatically buy on a regular basis, companies don't have to spend extra marketing money to target them. Also, subscribers, who pay periodically and are committed to a long-term association with the company, are usually less likely to churn. This reduces customer retention costs.

5.3 Popular Subscription Business Models

5.3.1 SaaS (Software-as-a-Service)

SaaS is a software delivery model where a software is centrally hosted on the cloud and licensed on a subscription basis.

The software is accessed via the internet, eliminating the need for installations or local infrastructure. Instead of purchasing software outright, customers pay a smaller, ongoing fee, reducing the initial investment.

5.3.2 Subscription Boxes

Subscription boxes are a form of subscription business model where customers receive well-chosen packages of products delivered to their doorstep on a recurring basis, typically monthly. Subscription boxes span various products from essentials, luxury goods, and meal kits, to even collector's items.

Brands use subscription boxes to introduce their products to new customers, often at discounted or bundled rates.

5.3.3 Publications & Newsletters

This subscription business model provides access to written or digital content such as news, reports, or entertainment in exchange for a recurring fee.

Subscribers receive content on a regular schedule, either digitally (email, apps, web) or in print. Newsletters often focus on specific industries, topics, or interests, delivering expert insights to a targeted audience.

5.3.4 Media Subscriptions

Media subscriptions, especially in video and music streaming, are a major form of the subscription business model where users pay a recurring fee for unlimited access to digital entertainment content.

Subscribers can stream content anytime, anywhere, no downloads or ownership required. Plans vary based on features like number of screens, audio/video quality, and ad-free experiences.

5.4 Metrics all Subscription Models track

5.4.1 Monthly Recurring Revenue (MRR)

MRR is the predictable recurring revenue earned from subscriptions in a particular month. MRR is an easy way to predict the future revenue and check on the business stability.

MRR is calculated by multiplying the average revenue per user (ARPU) by the total number of active subscribers in a month.

5.4.2 Annual Recurring Revenue (ARR)

ARR tells how much revenue subscription businesses can expect to generate from their subscribers every year. ARR is often looked at by investors to decide if the business has potential.

ARR is particularly useful for companies with annual billing cycles.

$$ARR = MRR \times 12$$

5.4.3 Average Revenue Per User (ARPU)

ARPU is the revenue generated from each paid subscription over a period, typically per month or year. ARPU helps to get a deeper understanding of the buyer persona and form a strategy for upgrades and add-ons.

$$ARPU = \text{Total Revenue} / \text{Number of Active Users}$$

5.4.4 Customer Lifetime Value (CLTV)

CLTV estimates the total revenue a business can expect from a single customer over the entire duration of their relationship. A guideline for businesses to stay profitable is to ensure CLTV is more than 3 times of the customer acquisition cost (CAC).

$$CLTV = ARPU \times \text{Average Customer Lifespan (in months or years)}$$

5.4.5 Customer Acquisition Cost (CAC)

CAC tells how much money a business spends in obtaining new customers. This can help companies analyse the efficiency of their marketing in bringing new customers aboard.

Tracking CAC allows businesses to optimize campaigns, reduce unnecessary spends on marketing, and improve profitability.

$$CAC = \text{Total Sales \& Marketing Expenses} / \text{Number of New Customers Acquired}$$

5.4.6 Churn Rate

Churn Rate measures the percentage of customers who cancel or do not renew their subscriptions over a specific period of time. It's a direct indicator of customer satisfaction and retention.

$$\text{Churn Rate} = (\text{Customers Lost During Period} / \text{Total Customers at Start of Period}) \times 100$$

5.5 Subscription Box Financial Model

The Subscription Box Financial Model evaluates the financial performance and profitability of subscription-based businesses. It demonstrates how revenue, costs, and customer-related metrics contribute to long-term sustainability. The key aspects of this model are:

1. **Revenue Model:** Subscription boxes generate predictable income through recurring monthly or annual fees. Key financial indicators include Monthly Recurring Revenue (MRR), Annual Recurring Revenue (ARR), and Average Revenue Per User (ARPU), which together form the basis of revenue forecasting.
2. **Customer Metrics and Retention:** Long-term profitability depends on maintaining a favorable ratio of Customer Lifetime Value (CLTV) to Customer Acquisition Cost (CAC). Monitoring the churn rate helps assess customer retention and predict future cash flows.
3. **Cost Structure and Profitability:** The model considers both fixed and variable costs, such as product sourcing, packaging, shipping, marketing, and overheads. Unit economics and break-even analysis are used to determine the subscription volume needed to reach profitability.
4. **Cost-Benefit and Scaling:** Subscription models benefit from predictable cash flows and lower retention costs. A cost-benefit evaluation ensures that the revenue generated and customer loyalty justify the associated expenses.
5. **Strategic Insights:** Profitability improves with better retention, structured pricing plans, and efficient acquisition strategies. Evaluating the Serviceable Obtainable Market (SOM) helps in planning realistic growth.

In summary, this model serves as a guide for revenue forecasting, scaling strategies, and financial planning in subscription-driven businesses.

5.6 Excel-Based Subscription Financial Model

To complement the conceptual understanding, an Excel-based financial model was developed to simulate 36 months of subscription box operations. This model demonstrates the connection between acquisition, revenue generation, cost structure, and profitability, providing a clear financial trajectory for the business.

1. **Unit Economics:** Revenue per box and the corresponding product, packaging, and shipping costs are used to calculate the contribution margin.

2. **Customer Acquisition and CAC:** Subscriber growth is modeled based on marketing efforts and Customer Acquisition Cost (CAC). Churn rate is applied to estimate the active subscriber base.
3. **Revenue and Growth:** Monthly Recurring Revenue (MRR) and Annualized Run Rate (ARR) are calculated to highlight the predictable nature of subscription revenue.
4. **Cost and Contribution Margin:** Both variable and fixed costs are included to determine monthly contribution margin and overall cash flow.
5. **Break-Even and Profitability:** The model identifies the point at which revenue exceeds total costs, marking the break-even and the start of profitability.

This hands-on model reinforced how acquisition metrics, cost structures, and retention efforts collectively drive sustainable subscription business growth.

Table 1: Key Insights from the Excel Subscription Model (Sample 6 Months)


Month	Revenue (INR)	Subscr. Share	Transac. Share	CAC (INR)	LTV:CAC
M1	43,164	32%	68%	100	1.09
M2	97,571	39%	61%	100	1.09
M3	1,60,517	43%	57%	100	1.09
M4	2,80,680	51%	49%	75	1.51
M5	3,95,679	56%	44%	75	1.51
M6	5,15,333	59%	41%	75	1.51

This table summarizes how revenue growth, the shift toward subscription revenue, and improvements in CAC collectively enhance the LTV:CAC ratio, which is critical for scaling a subscription business effectively. Link to the excel: [Sheet](#)

6 Cost-Benefit Analysis

6.1 What is Cost-Benefit Analysis ?

Cost-Benefit Analysis (CBA) is a decision-making tool used to evaluate the financial, social, and environmental costs and benefits of a project or investment. The goal is to determine whether the benefits outweigh the costs, and by how much.




The **COSTS** of a project can include :

- Capital Costs (e.g. equipment, technology, etc.)
- Development and Implementation Costs (e.g. labor, services, etc.)
- Ongoing and Maintenance Costs (e.g. licensing, upgrades, repairs ,etc.)

Some typical **BENEFITS** of a project include :

- **Cost Savings Through Reduced Operational Expenditures**
These include minimizing key operational expenses, reduced labor requirements, lower dependency on external services, decreased downtime, optimized use of supplies, and decreased product costs.
- **Increased Revenue and Enhanced Customer Value**
The project is anticipated to drive higher revenue by attracting more customers, increasing customer loyalty and advocacy, and encouraging greater spending per customer.
- **Intangible Benefits: Improved Morale and Brand Perception**
The project is expected to generate soft benefits such as enhanced employee morale and a more positive consumer perception of the brand.



6.2 Time Factor in Cost-Benefit Analysis

Time is a large factor in doing a cost-benefit analysis, in particular, how the costs and benefits evolve. Do they last five years, or ten years? Do they improve or degrade over time? The value of an investment is based on the return of the investment over time. Analysts must consider not just the total amount of costs and benefits, but also when they occur.

Another way to calculate cost-benefit is to calculate the payback period. Payback period is the time taken by an investment to break even. For example, if implementation costs equal 10,000 , and the annual benefits minus ongoing costs equal 5,000, then there is a 2-year payback on the investment.

6.3 Best Practices in Cost-Benefit Analysis

Identify the Right Cost and Benefit Categories : Even the most detailed analysis can lead to misleading conclusions if it overlooks the major cost and benefit drivers. It's essential to start by clearly defining and prioritizing the key categories that matter most.

Ground Your Assumptions in Real-World Insights : While rough estimates are sometimes unavoidable, it's critical to anchor your cost and benefit projections in actual data whenever possible. Alternatively, surveys, interviews, or focus groups can provide valuable input straight from the affected customers.

Quantify Wherever Possible ; Use Ranges or Scales When You Can't : Whenever possible, tie each cost and benefit to a specific dollar value. But when precise quantification isn't feasible, assign a relative magnitude such as "low," "moderate," or "high" to reflect its potential impact.


Continuously Estimate Rough Cost-Benefit Figures : It is good practice to chart out the cost and benefit categories and estimate the order of magnitude of the categories tied to an idea or potential initiative. It often leads to great debates that build out the idea and helps orient the relative priority of other ideas and initiatives.

7 How to Calculate the Break-Even Point

7.1 What is the Break-Even Point ?

The break-even point (BEP) is the point at which **total costs = total revenue**, meaning there is no profit and no loss. It's a critical financial metric used to determine how much you need to sell (in units or in revenue) to cover your costs.

It tells when a company, or one of its products, will start to be profitable. If a business's revenue is below the break-even point, then the company is operating at a loss. If it's above, then it's operating at a profit.



7.2 How to Calculate the Break-Even Point ?

7.2.1 In Units

$$\text{Break-even point (units)} = \frac{\text{Fixed Costs}}{\text{Selling Price per Unit} - \text{Variable Cost per Unit}}$$

Fixed Costs - Fixed costs are ones that typically do not change, or change only slightly. Examples of fixed costs for a business are monthly utility expenses and rent.

Selling Price per Unit - This is how much a company is going to charge consumers for just one of the products that the calculation is being done for.

Variable Costs per Unit - Variable costs are the costs that often fluctuate, and are typically a company's largest expense.

7.2.2 In Revenue

$$\text{Break-even point (revenue)} = \frac{\text{Fixed Costs}}{\text{Contribution Margin Ratio}}$$

$$\text{Contribution Margin} = \text{Selling Price} - \text{Variable Cost}$$

$$\text{Contribution Margin Ratio} = \frac{\text{Contribution Margin}}{\text{Selling Price}}$$

8 Metrics and KPI

This section outlines key metrics and KPIs to monitor the effectiveness of the KahaniBox platform across different stages of the user journey and business objectives.

8.1 Customer Retention Rate (CRR)

Formula: $\text{CRR} = \frac{E-N}{S} \times 100$

Where:

- E = number of customers at the end of the period
- N = number of new customers acquired during the period
- S = number of customers at the start of the period

A higher CRR indicates better user experience, content engagement, and platform loyalty.

8.2 Net Promoter Score (NPS)

Measures user satisfaction and willingness to recommend the platform.

Formula: $\text{NPS} = \% \text{Promoters} - \% \text{Detractors}$

Useful for capturing qualitative feedback and understanding user sentiment.

8.3 Churn Rate

Percentage of users who stop using the platform over a given time period.

Formula: $\text{Churn Rate} = \frac{\text{Users Lost}}{\text{Total Users at Start}} \times 100$

Helps identify retention problems and optimize content and experience.

8.4 Monthly Active Users (MAU)

Counts unique users who interacted with the platform in the last 30 days. Indicates overall product usage and reach.

8.5 Daily Active Users (DAU)

Tracks daily engagement. High DAU indicates consistent user interest and habit formation.

8.6 Activation Rate

Measures how many new users perform a key activation action (e.g., watching a story, saving, or sharing) within their first session or day.

Formula: $\text{Activation Rate} = \frac{\text{Users Who Activated}}{\text{New Users}} \times 100$

Essential for evaluating the effectiveness of onboarding and story discovery.

8.7 Average Session Length

Indicates how long users are spending per session. A high session length shows engaging storytelling and platform design.

8.8 Customer Acquisition Cost (CAC)

Formula: $\text{CAC} = \frac{\text{Total Marketing and Sales Cost}}{\text{Number of New Customers Acquired}}$

Important for measuring efficiency of marketing strategies.

8.9 Lifetime Value (LTV)

Estimates the revenue a customer will generate during their lifetime with the platform.

Formula: $\text{LTV} = \text{ARPU} \times \text{Gross Margin} \times \text{Customer Lifetime}$

8.10 LTV : CAC Ratio

Assesses the profitability of acquisition strategies. Ideal ratio: >**3:1**.

8.11 Story Completion Rate

Measures how many users complete a story once started.

Formula: $\text{Completion Rate} = \frac{\text{Stories Completed}}{\text{Stories Started}} \times 100$

Indicates quality and appeal of content.

8.12 Engagement Rate

Tracks how often users perform interactive actions (like, share, comment, bookmark).

Formula: $\text{Engagement Rate} = \frac{\text{Total Engagement Actions}}{\text{Total Users}} \times 100$

8.13 Virality Coefficient

Measures how many new users are brought in by existing users through referrals or shares.

Formula: $\text{K-Factor} = \text{Average Invitations per User} \times \text{Conversion Rate}$

$K > 1$ implies exponential growth potential.

8.14 Revenue Metrics

- ARPU (Average Revenue Per User)
- MRR (Monthly Recurring Revenue)
- ARR (Annual Recurring Revenue)
- Revenue Churn Rate –

8.15 Funnel Conversion Rates

- Visitor to Sign-Up Rate
- Sign-Up to Activation Rate
- Activation to Retention Rate
- Retention to Monetization Rate

Helps optimize weak stages in the user journey.

8.16 Bounce Rate

Percentage of users who leave the app without taking any meaningful action. Should be minimized.

8.17 DAU/MAU Ratio

Measures user stickiness. **Formula:** $\frac{\text{DAU}}{\text{MAU}} \times 100$

A good product often has a ratio around 20–30

DATA SCIENCE

9 Introduction to Python Libraries

What is Python?

Python is a high-level, interpreted, general-purpose programming language. It supports multiple programming paradigms, including object-oriented, imperative, and functional programming. Python libraries pandas, NumPy, Matplotlib, and Seaborn were utilized to handle, analyze, and visualize data efficiently. These libraries were integral in transforming raw data into meaningful insights.

9.1 NumPy: Numerical Python

9.1.1 Overview

NumPy (Numerical Python) is the fundamental package for numerical computing in Python. It provides support for large, multi-dimensional arrays and matrices, along with a collection of high-level mathematical functions to operate on these arrays. NumPy is the backbone of many other scientific computing libraries in Python, including Pandas and Matplotlib.

9.1.2 Key Features and Functionalities

- **ndarray Object:** The core of NumPy is the ‘ndarray’ object, an efficient multi-dimensional array. This array is homogeneous, meaning all elements must be of the same data type, which significantly boosts performance compared to Python lists for numerical operations.
- **Broadcasting:** NumPy’s broadcasting mechanism allows for efficient operations between arrays of different shapes.
- **Mathematical Functions:** A comprehensive set of mathematical functions for linear algebra, Fourier transform, random number generation, etc.
- **Performance:** NumPy operations are implemented in C and Fortran, making them highly optimized for speed.
- **Memory Efficiency:** ‘ndarray’ objects are more memory-efficient than Python lists for numerical data.

9.1.3 Example Usage

NumPy is used for tasks like:

- Creating and manipulating arrays.
- Performing element-wise operations.
- Linear algebra operations (e.g., dot product, matrix inversion).
- Statistical computations (mean, median, standard deviation).

9.2 Pandas: Data Analysis Library

9.2.1 Overview

Pandas is a fast, powerful, flexible, and easy-to-use open-source data analysis and manipulation library for Python. It is built on top of NumPy and provides data structures and functions needed to work with structured data seamlessly. Pandas is widely used in data cleaning, transformation, analysis, and exploration.

9.2.2 Key Data Structures

- **Series:** A one-dimensional labeled array capable of holding any data type (integers, strings, floats, Python objects, etc.). It is similar to a column in a spreadsheet or a SQL table.
- **DataFrame:** A two-dimensional labeled data structure with columns of potentially different types. It is similar to a spreadsheet, a SQL table, or a dictionary of Series objects. DataFrames are the most commonly used Pandas object.

9.2.3 Key Features and Functionalities

- **Data Alignment:** Handles missing data and automatically aligns data based on labels.
- **Data Loading/Saving:** Easy reading and writing of data from various file formats (CSV, Excel, SQL databases, JSON).
- **Data Cleaning:** Powerful tools for handling missing data, duplicate data, and data type conversions.
- **Data Manipulation:** Filtering, sorting, grouping, merging, and reshaping data.
- **Time Series Functionality:** Robust tools for working with time-series data.

9.2.4 Example Usage

Pandas is essential for:

- Reading data from CSV files into DataFrames.
- Cleaning and preparing messy datasets.
- Performing statistical analysis on tabular data.
- Aggregating and summarizing data.

9.3 Matplotlib: Plotting Library

9.3.1 Overview

Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python. It provides an object-oriented API for embedding plots into applications using general-purpose GUI toolkits like Tkinter, wxPython, Qt, or GTK. Matplotlib is widely used for generating high-quality figures for publications, presentations, and reports.

9.3.2 Example Usage

Matplotlib is used for:

- Visualizing trends in data with line plots.
- Showing relationships between variables with scatter plots.
- Displaying distributions with histograms.
- Creating professional-looking charts and graphs.

10 Machine Learning

10.1 Introduction to ML

Machine Learning (ML) marks a paradigm shift in computing—from traditional programming, where humans write explicit rules, to a data-driven approach, where humans provide data and the machine learns patterns and predictions from it.

The primary goal of machine learning is to enable machines to learn from data in order to make predictions, detect patterns, or take actions without being explicitly programmed for specific tasks, this can include

Classification involves assigning inputs to discrete categories. A common example is recognizing handwritten digits using pixel grayscale values.

Regression deals with predicting continuous values, such as forecasting stock market prices using statistical patterns.

Instead of manually coding solutions, ML systems are trained using data and models. The machine "learns" from the data to make future predictions or decisions.

There are several types of learning in ML:

1. Supervised Learning: The model is trained on labeled data (input-output pairs) to learn how to predict outputs for new inputs.
2. Unsupervised Learning: The model identifies patterns or structures in data without labeled outputs.
3. Reinforcement Learning: The model learns through trial and error, receiving rewards or penalties based on its actions.

In this discussion, we focus on supervised learning models, specifically linear regression, which is a foundational technique used to model relationships between variables and make predictions.

10.2 Mean based model

Say we have data points x_i, y_i , each of these points is colored red and blue, we have a data set with maps x_i to y_i , now given a point and x, y we aim to create a model which predicts its color. We make our model using the mean of the norms of the red points μ_r and the mean of the norms of the blue points as μ_b . we calculate distance of a point from this mean

$$||\mu_r - x||^2 = |\mu_r|^2 + |x|^2 - 2\langle\mu_r, x\rangle \quad (1)$$

$$||\mu_b - x||^2 = |\mu_b|^2 + |x|^2 - 2\langle\mu_b, x\rangle \quad (2)$$

$$f(x) = ||\mu_r - x||^2 - ||\mu_b - x||^2 = 2\langle\mu_b - \mu_r, x\rangle + ||\mu_r||^2 - ||\mu_b||^2 \quad (3)$$

$$f(x) = \langle w, x \rangle + b \quad (4)$$

We have successfully reduced to a linear model. we predict red if $f(x) > 0$ else we predict blue. This model show how a ML model might look like, but this model has many flaws.

1. Assumes each class is well represented by a single mean vector.
2. fails if classes have overlapping regions, making linear separation hard.
3. fails if each class has multiple distant clusters (i.e., is multi-modal).

We also learn that even though this model fails in general but it may work on compact, unimodal and linearly separable. Hence, model selection is correlated with data available.

11 Linear Regression

11.1 Introduction

Linear Regression is a foundational supervised learning technique used to predict a continuous target variable based on one or more input features. It is widely adopted due to its simplicity, transparency, and effectiveness in identifying trends and relationships within numerical data. From economics and marketing to healthcare and urban planning, linear regression remains a go-to tool for data-driven forecasting and decision-making. The most common method for estimating the parameters of a linear regression model is Ordinary Least Squares (OLS), which finds the line that best fits the data by minimizing the sum of squared errors.

11.2 Simple Linear Regression

Simple linear regression models the relationship between one independent variable X and one dependent variable Y :

$$Y = \beta_0 + \beta_1 X + \epsilon$$

- β_0 : Intercept
- β_1 : Slope (change in Y for a unit change in X)
- ϵ : Error term (residual)

11.3 Multiple Linear Regression

With multiple independent variables, the model becomes:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_n X_n + \epsilon$$

11.4 Assumptions

- **Linearity:** The relationship between independent and dependent variables is linear.
- **Independence:** Residuals (errors) are independent.
- **Homoscedasticity:** The variance of residuals is constant across all levels of the input variables. Variance here refers to the average squared deviation from the mean. A non-constant spread of residuals indicates violation.
- **Normality of Errors:** The residuals should be normally distributed, especially when constructing confidence intervals or conducting hypothesis testing.
- **No Multicollinearity:** In multiple regression, independent variables should not be highly correlated.

11.5 Mathematical Formulation

The parameters $\beta_0, \beta_1, \dots, \beta_n$ are estimated using the Ordinary Least Squares (OLS) method:

$$\min \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

In matrix form, the closed-form solution is:

$$\hat{\beta} = (X^T X)^{-1} X^T Y$$

11.6 Example and Interpretation

Consider a real-world scenario of predicting house prices using various features like:

- Size of the house (in square feet)
- Number of bedrooms and bathrooms
- Age of the house
- Distance from city center
- Whether the house has a garage
- Locality or neighborhood score

A linear regression model can quantify how each of these features influences the final price. For example:

- **House Size:** Typically, prices increase with area. The model estimates how much price changes per additional square foot.

- **Rooms:** More rooms add value, but returns may diminish.
- **Age:** Newer homes may be valued more, unless older homes have been renovated.
- **Garage:** A yes/no feature that could boost value.
- **Locality:** Higher-rated neighborhoods may significantly increase pricing.

Once trained, the model can be used to predict the price of a new property using its features. In addition to prediction, the regression model helps stakeholders identify which variables are most impactful in determining value and discover pricing trends across the dataset.

11.7 Applications

- **Finance:** Stock price forecasting, credit scoring
- **Healthcare:** Disease progression and treatment cost estimation
- **Marketing:** Customer lifetime value prediction, price optimization
- **Education:** Predicting student performance based on attendance and engagement
- **Operations:** Inventory demand forecasting

11.8 Linear regression and gradient descent

In problems we have X and Y input and output variables, X and Y can be in R^n . Example : say we want to predict prices of housing based on no. of bedrooms and total area then X is in R^2 and Y in R . Our aim is to find $h : X \rightarrow Y$

Our hypothesis is $h_\theta(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 \dots$ here $\theta = [\theta_0, \theta_1, \theta_2 \dots]$ is parameter vector and $x = [x_0, x_1, x_2 \dots]$ is feature vector. we take $x_0 = 1$. $h(x) = \theta^T x$. We aim to find θ such that our error function called cost function

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 \quad (5)$$

We aim to minimize $J(\theta)$ by starting with an initial guess for θ iteratively update θ to reduce $J(\theta)$. we define update rule as

$$\theta_j = \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j} \quad (6)$$

for $j = 0, 1, \dots, n$

here α is called the learning rate and the update moves in the direction of steepest descent. defining $x_j = \frac{\partial (h_\theta(x) - y)}{\partial \theta_j}$ we get

$$\theta_j = \theta_j + \alpha (y - h_\theta(x)) x_j \quad (7)$$

apply on full data set $\{(x^i, y^i)\}_{i=1}^m$ we now get batch gradient descent as

$$\theta_j = \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_\theta(x^{(i)})) x_j^{(i)} \quad (8)$$

This uses entire dataset on every step which makes it computationally expensive, we always reach the global minimum of θ if α is small enough.

Stochastic Gradient descent is another model which does not use the whole dataset for updating, rather it updates one by one, this helps if the data set is very large, also it can start learning immediately and escape shallow minima. Its algorithm is as follows.

for $i = 1$ to m : (for all j)

$$\theta_j = \theta_j + \alpha (y^{(i)} - h_\theta(x^{(i)})) x_j^{(i)} \quad (9)$$

Its drawbacks include

1. Does not converge exactly
2. θ may keep oscillating around the minimum
3. Noisy updates due to individual examples.

11.9 Non iterative approach

Here we aim to minimize $J(\theta)$ by directly calculating where its gradient is zero.

We write matrix X as **Design Matrix** X :

$$X = \begin{bmatrix} - & - & - & (x^{(1)})^T & - & - & - \\ - & - & - & (x^{(2)})^T & - & - & - \\ & & & \vdots & & & \\ - & - & - & (x^{(m)})^T & - & - & - \end{bmatrix} \in \mathbb{R}^{m \times (n+1)}$$

Output Vector y :

$$\vec{y} = \begin{bmatrix} y^{(1)} \\ y^{(2)} \\ \vdots \\ y^{(m)} \end{bmatrix} \in \mathbb{R}^m$$

Our hypothesis remains same $h_\theta(X) = X\theta$ where error vector is $X\theta - \vec{y}$ we minimize

$$J(\theta) = \frac{1}{2}(X\theta - \vec{y})^T(X\theta - \vec{y}) \quad (10)$$

$$= \frac{1}{2}[\theta^T X^T X \theta - 2\vec{y}^T X \theta + \vec{y}^T \vec{y}] \quad (11)$$

taking gradient with respect to θ and equating to zero for minimizing $J(\theta)$

$$\begin{aligned} \nabla_\theta(J(\theta)) &= X^T X \theta - X^T \vec{y} \\ X^T X \theta - X^T \vec{y} &= 0 \\ X^T X \theta &= X^T \vec{y} \\ \theta &= (X^T X)^{-1} X^T \vec{y} \end{aligned} \quad (12)$$

equation (12) is also called *normal equation* and is the required solution.

12 Unsupervised Learning

Unsupervised learning is a type of machine learning where the algorithm works with data that has no predefined labels or outcomes. Unlike supervised learning, where models are trained using input-output pairs, unsupervised learning identifies hidden patterns, relationships, or structures in the data without any prior guidance. It is particularly useful for tasks such as clustering, dimensionality reduction, and anomaly detection. Since the data is unlabeled, the model must analyze and interpret the input independently, making it a powerful tool for exploring complex datasets and uncovering insights that might not be immediately apparent.

12.1 How Unsupervised Learning Works?

- **Input of Unlabeled Data:** The model is fed raw data without any predefined labels or categories.
- **No Prior Knowledge:** The algorithm identifies that it must learn patterns without external guidance.
- **Pattern Discovery:** Techniques like **clustering** and **dimensionality reduction** are used to find structure in the data.
- **Formation of Groups:** Data points are grouped based on similarity or feature relationships.
- **Meaningful Output:** The result is organized clusters or patterns that reveal insights.
- **Key Idea:** The model learns hidden structures in data without human supervision.

12.2 Clustering in Unsupervised Learning

- **Clustering** is a core technique in unsupervised learning that involves grouping data points based on similarity or shared features.
- It allows the algorithm to **discover hidden patterns or natural groupings** in datasets without the need for labeled data.
- Clustering helps reveal structure in the data, such as grouping similar customers, documents, or pixels.
- It is commonly used in applications like:
 - **Customer segmentation** (e.g., targeting marketing strategies),
 - **Image compression** (by grouping similar pixel values),
 - **Anomaly detection** (identifying outliers or unusual behavior).
- The goal is to maximize similarity within clusters and minimize similarity between clusters.

12.3 How Clustering Reveals Structure

Clustering reveals hidden structure in unlabeled data by grouping similar data points based on their features. It identifies natural patterns that aren't immediately visible in raw data.

By measuring similarity (e.g., distance between points), clustering algorithms form distinct groups where points within a cluster are more alike than those in different clusters. This helps uncover meaningful segments, trends, or associations, making complex data easier to interpret and analyze.

Example: Clustering Customers by Salary and Spending Score

This example illustrates how clustering can reveal meaningful structure in unlabeled data using two features: **Salary** and **Spending Score**.

- In the **left plot**, raw data is shown without any labels. The data points appear randomly scattered, with no clear grouping or pattern.
- In the **right plot**, after applying a clustering algorithm (e.g., K-Means), the same data is grouped into distinct clusters based on similarities in income and spending behavior.

Each color in the clustered output represents a group of customers with similar financial patterns:

- Low Salary, Low Spending
- Low Salary, High Spending
- High Salary, Low Spending
- High Salary, High Spending


12.4 K-Means Clustering

12.4.1 How K-Means Clustering Works


- The algorithm begins by **randomly selecting a few centroids** (central points in the data space).
- Each data point is **assigned to the nearest centroid**, forming the initial clusters.
- Once all points are assigned, the **centroids are updated** as the **mean of all points in their respective clusters**.
- This process of **assignment and centroid update is repeated** iteratively.
- The algorithm stops when **centroids stabilize**, meaning they no longer change significantly between iterations.
- The final output is a set of **well-defined clusters** where similar data points are grouped together based on feature similarity.

The goal of K-Means is to divide data into k distinct clusters such that data points within the same cluster are more similar to each other than to those in other clusters. This ensures that each cluster captures a meaningful group based on similarity of characteristics.

Raw Input Data



Grouped Clusters (After Clustering)



12.4.2 Intuition Behind K-Means Clustering

- **Initialization** Randomly select k initial points as cluster centroids (also called means).
- **Assignment Step** Assign each data point to the nearest centroid based on distance (typically Euclidean).
- **Update Step** For each cluster, recalculate the centroid as the average position of all points assigned to it.
- **Iteration** Repeat the assignment and update steps for a fixed number of iterations or until the centroids no longer move significantly.

By the end of the process, similar data points are grouped into distinct clusters, each represented by a final centroid. This results in clear segmentation of the dataset based on natural similarities.

12.4.3 WCSS and the Elbow Method

What is WCSS?

The Within-Cluster Sum of Squares (WCSS) measures how close data points are to the centroids of their respective clusters. It reflects the compactness of the clusters — lower values indicate that data points are closely grouped around the centroid.

WCSS is calculated using the formula:

$$\text{WCSS} = \sum_{i=1}^k \sum_{x \in C_i} \|x - \mu_i\|^2$$

where:

- k is the number of clusters,
- $x \in C_i$ represents a data point in cluster C_i ,
- μ_i is the centroid of cluster C_i ,
- $\|x - \mu_i\|^2$ is the squared Euclidean distance between the point and the centroid.

WCSS is commonly used to evaluate clustering performance and is a key component in the Elbow Method for selecting the optimal number of clusters.

How the Elbow Method Works

- Calculate **WCSS (Within-Cluster Sum of Squares)** for different values of k (number of clusters).
- Plot a graph with **k on the x-axis** and **WCSS on the y-axis**.
- Observe the shape of the curve as k increases.
- WCSS will **decrease with increasing k** , but the rate of decrease slows down after a certain point.
- The point where the curve **starts to flatten** is called the "**elbow point**".
- This elbow represents the **optimal number of clusters**, where adding more clusters yields minimal improvement.

12.4.4 Random Initialization Trap in K-Means

Problem: Random Initialization in K-Means

- K-Means starts by randomly selecting initial centroids.
- Poor initialization can lead to:
 - Convergence to suboptimal (local) minima
 - Inconsistent results across different runs
 - Uneven or empty clusters

Solution: Improving Initialization

- Use K-Means++, which selects better initial centroids by spreading them out, improving stability and convergence.
- Run K-Means multiple times with different initializations and choose the result with the lowest WCSS for better clustering performance.

k-means++ is an improved initialization method for K-Means clustering. Instead of choosing centroids completely at random, it selects the initial centroids in a way that spreads them out across the data space. This reduces the chances of poor clustering results and helps the algorithm converge faster and more reliably. By starting with well-separated centroids, k-means++ typically leads to better cluster quality and more consistent outcomes compared to standard random initialization.

12.5 Hierarchical Clustering

12.5.1 What is Hierarchical Clustering?

- A clustering method that builds a tree-like structure (dendrogram) to represent data grouping.
- Starts with each data point as its own cluster.
- At each step, the two closest clusters are merged based on a distance metric (e.g., Euclidean distance).
- This merging continues until all data points are combined into one large cluster.
- The resulting hierarchy can be visualized using a dendrogram.
- By cutting the dendrogram at a specific level, a desired number of clusters can be formed.

Example: Clustering Fruits by Weight

Four fruits: Apple (100g), Banana (120g), Cherry (50g), Grape (30g)

Step 1: Cherry and Grape are grouped first (lightest)

Step 2: Apple and Banana are grouped (heavier pair)

Step 3: All fruits are merged into a single cluster

This process forms a hierarchy, which can be visualized using a dendrogram.

12.5.2 Dendrogram: A Tree for Clusters

What is a Dendrogram?

- A **dendrogram** is a **tree-like diagram** that visualizes the merging process in hierarchical clustering.

-
- It acts like a **family tree for clusters**, showing how data points or groups are combined step by step.
 - At the **bottom level**, each data point starts as its own cluster.
 - As you move **upward**, similar points or clusters are **merged into larger groups**.
 - The **height at which two clusters merge** reflects their similarity —
 - **Lower merge point = more similar** groups
 - **Higher merge point = less similar** groups
 - By drawing a **horizontal cut** at a certain height, you can choose the number of clusters.

Dendrograms help visualize the entire clustering process step by step and assist in deciding the number of clusters.

12.5.3 Types of Hierarchical Clustering

Agglomerative Clustering (Bottom-Up): Start with individual data points. Merge closest pairs until one big cluster is formed.

Divisive Clustering (Top-Down): Start with one large cluster. Recursively split clusters into smaller ones

12.5.4 Inter-Cluster Distance (Linkage Methods)

When merging two clusters, we must decide how to measure the distance between them. Different methods define "distance" differently

- Minimum Distance (Single Linkage):
 - Distance between the closest pair of points in two clusters.
 - Tends to form long, chain-like clusters.
- Maximum Distance (Complete Linkage):
 - Distance between the farthest pair of points in two clusters.
 - Tends to create compact, spherical clusters.
- Group Average (Average Linkage):
 - Average distance between all pairs of points across the two clusters.
 - A balance between single and complete linkage.
- Ward's Method:
 - Merges clusters that result in the smallest increase in total within-cluster variance.
 - Tends to produce equally sized, spherical clusters

12.5.5 How to Select the Number of Clusters?

In hierarchical clustering, we don't pre-define the number of clusters. Instead, we can choose it after building the dendrogram by cutting it at a desired level.

Approaches to Select Number of Clusters:

Using Dendrogram:

- Visually inspect the dendrogram.
- Look for a large vertical gap (height) between merges.
- Draw a horizontal cut that intersects the largest vertical distance.
- The number of vertical lines it cuts gives the number of clusters

13 Principal Component Analysis (PCA)

13.1 What is PCA?

Principal Component Analysis (PCA) is an unsupervised statistical technique used for dimensionality reduction. It transforms high-dimensional data into a lower-dimensional form by identifying the directions (principal components) along which the variation in the data is maximal.

Example: Suppose you have a dataset of students with two correlated features: math score and physics score. PCA finds a new axis (principal component) along which the scores vary the most—this might represent overall science ability. Instead of analyzing both scores separately, PCA allows you to analyze just this new component, reducing dimensions while keeping most of the information.

13.2 Benefits of PCA

PCA is widely used in data analysis and machine learning for several reasons:

- **Dimensionality Reduction:** Reduces the number of features while preserving as much information as possible.
- **Noise Reduction:** By discarding components with low variance (assumed to be noise), PCA can help denoise the data.
- **Visualization:** PCA enables visualization of high-dimensional data in 2D or 3D space using the top principal components.
- **Improved Performance:** Reducing input dimensions may improve performance and reduce overfitting in machine learning models.

13.3 Why Variance is Important?

Variance measures how much the data is spread out. In PCA, the goal is to retain directions where the data varies the most — because variance indicates the amount of information or structure present.

If a direction in feature space has high variance, it means that data points are more dispersed in that direction, potentially holding important patterns.

PCA selects principal components such that:

- The first principal component maximizes the variance.
- Subsequent components maximize variance orthogonal to the prior components.

Hence, variance directly guides the ranking and selection of principal components.

13.4 Covariance

Covariance is a statistical measure of how two variables change together. In PCA, understanding how features covary helps us find directions where data shows meaningful patterns.

Covariance Formula

Given two variables X and Y , their covariance is defined as:

$$\text{Cov}(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$$

Covariance Matrix

For a dataset with d features, the covariance matrix $\mathbf{C}$ is a $d \times d$ matrix where each element C_{ij} is the covariance between feature i and feature j :

$$\mathbf{C} = \frac{1}{n-1} (\mathbf{X} - \bar{\mathbf{X}})^T (\mathbf{X} - \bar{\mathbf{X}})$$

This matrix is symmetric, and its eigenvectors define the principal components.

Significance in PCA

- Eigenvectors of the covariance matrix define the directions of maximum variance (principal components).
- Eigenvalues indicate the magnitude of variance along each component.
- PCA uses this matrix to transform data into uncorrelated axes ordered by importance.

13.5 Geometric Interpretation of PCA

Geometrically, PCA rotates the original coordinate axes to align with the directions of greatest variability in the data.

- The principal components are orthogonal vectors that form a new basis.
- These vectors point in the directions where the data is most spread out.
- Each data point is projected onto this new basis, reducing dimensionality.




Figure 1: Geometric view: PCA aligns data to new axes of maximum variance

- Original axes: X and Y represent the original coordinate system.
- The data (green points) is scattered in an elongated shape, showing correlation.
- PC_1 : Direction of maximum variance (first principal component).
- PC_2 : Orthogonal to PC_1 , captures remaining variance.
- PCA rotates the axes to align with PC_1 and PC_2 , simplifying structure.

14 Decision Tree

14.1 What is a Decision Tree?

A **Decision Tree** is a **supervised learning algorithm** used for:

- **Classification**: Predicting discrete class labels.
- **Regression**: Predicting continuous values.

Key Components:


- **Root Node**: Represents the entire dataset.
- **Branches**: Represent attribute values.
- **Internal Nodes**: Represent decisions or attribute tests.
- **Leaf Nodes**: Represent final outcomes or predictions.

Decision Trees function like flowcharts for decision making.

14.2 How Does a Decision Tree Work?

- Splits data based on **feature values** to form pure subsets.
- Internal nodes make decisions based on feature conditions.
- Leaf nodes contain class labels or predicted values.

14.3 Decision Tree Structure Overview



14.4 Splitting Criteria in Classification Trees

a. Gini Index

$$I_G = 1 - \sum_{j=1}^c p_j^2$$

where p_j is the proportion of samples of class j .

b. Entropy

$$I_H = - \sum_{j=1}^c p_j \log_2(p_j)$$

Measures impurity; 0 if all items belong to one class.

14.5 Information Gain

- Measures reduction in entropy after a split.
- Helps identify the best attribute for splitting.
- Higher Information Gain implies a better split.

14.5.1 Formula:

$$Gain(S, A) = Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} \cdot Entropy(S_v)$$

Where:

- S : Full dataset
- A : Attribute
- S_v : Subset of S where $A = v$

14.6 Decision Trees for Regression

Decision Tree Regressors predict continuous values and use **variance reduction** or **mean squared error** for splitting.

7. Pruning in Decision Trees

Problem:

Fully grown trees often overfit the training data, capturing noise and reducing generalization.

Solution:

Use **pruning techniques** to simplify the tree and improve generalization.

Types of Pruning:

1. **Pre-Pruning (Early Stopping)**: Stop growing the tree when a condition is met.
 - *Common Stopping Criteria*:
 - Maximum depth
 - Minimum samples to split
 - Minimum gain in information or variance reduction
 - Maximum number of leaves
 - *Pros*: Fast, low memory, avoids deep trees
 - *Cons*: Risk of underfitting, may miss key splits
2. **Post-Pruning (Prune After Full Growth)**: Grow the full tree, then prune subtrees.
 - *Methods*:
 - **Reduced Error Pruning**: Remove nodes that don't worsen validation error.
 - **Cost-Complexity Pruning (CART)**:

$$R_\alpha(T) = R(T) + \alpha \cdot |T|$$

Where:

- * $R(T)$: Error of the subtree
- * $|T|$: Number of terminal nodes
- * α : Regularization parameter
- *Pros*: Better generalization, often more accurate
- *Cons*: Slower, needs validation set

Why Prune?

- Avoid overfitting
- Improve interpretability
- Reduce model complexity
- Enhance generalization